

NixOS as a declarative and synchronised solution to embedded security challenges and system administration problems for multiple embedded devices

TURUN YLIOPISTO
Tietotekniikan laitos
TkK-tutkielma
Labran nimi?
December 2024
Eino Korte

TURUN YLIOPISTO

Tietotekniikan laitos

EINO KORTE: NixOS as a declarative and synchronised solution to embedded security challenges and system administration problems for multiple embedded devices

TkK-tutkielma, A-6 s., 6 liites.

Labran nimi?

December 2024

Embedded devices are an integral part of our everyday lives; household machines, automobiles, and thermal sensors make use of embedded devices at all times. They are subject to the global, developing world's security problems. This thesis focuses on those found in public information screens. Embedded devices are particularly vulnerable to security problems as they have challenges receiving constant, reliable updates. This thesis' focal point is maintaining, updating and upgrading embedded devices. A proposed architecture solution is provided with example snippets to cover most of common security issues found in similar setups. The architecture with it's content is then evaluated through common systems security methodologies. Used methodologies are compared to more common methodologies. The central theme of this thesis is NixOS, which is a Linux distribution that forms itself from a set of configuration files, supporting features like atomic rollbacks and reliable dependency handling. Most definitive academic sources in this particular subject by Eelco Dostra are used extensively, as well as papers regarding both embedded security and systems security in general. Ideas for further study are presented, as security problems may arise in our everyday lives due to the more mainstream paradigms and could be avoided with the use of declarative ones.

Asiasanat: tähän, lista, avainsanoista

TURUN YLIOPISTO

Tietotekniikan laitos

EINO KORTE: NixOS as a declarative and synchronised solution to embedded security challenges and system administration problems for multiple embedded devices

TkK-tutkielma, A-6 s., 6 liites.

Labran nimi?

December 2024

Second abstract in english (in case the document main language is not english)

Asiasanat: nix, nixos, declarative, security, pomdp, queries

Contents

1	Introduction	1
1.1	Research methodologies and questions	3
1.1.1	Literature review	4
1.1.2	Research questions	5
1.1.3	Data collection and analysis	6
2	Declarative vs. Imperative systems	7
2.1	Imperative systems	8
2.1.1	Debian/Apt	9
2.1.2	Gentoo/Portage	10
2.2	Declarative systems	11
2.2.1	Non-declarative components	13
2.2.2	Home manager and flakes	15
2.2.3	Ease of updates	15
3	Embedded system security	17
3.1	Common embedded pitfalls	19
3.2	NixOS as an embedded solution	20
3.3	Imperative and declarative systems from CIA-triad approach	21
4	Proposed architecture solution	24

4.1	Server	24
4.2	Client	26
4.2.1	MQTT-client	26
4.2.2	Weston/XWayland	27
4.3	Test device	30
4.4	Installing new devices	30
5	Security standpoints	34
5.1	A brief history of security metrics	34
5.2	Choosing security metrics	35
5.3	Measuring security	36
5.4	Methodologies in comparison	38
5.5	Quantitative metrics	39
5.5.1	QuERIES	39
5.5.2	QuERIES as a central methodology	40
5.5.3	Partially observable Markov decision process	42
6	Analysis	44
6.1	Modeling the problem and quantifying the models	44
6.2	Modeling the possible attacks	45
6.3	Using the results	48
6.3.1	Red team implications	48
6.3.2	Blue team implications	48
7	Further research	49
7.1	Further research questions	49
8	Conclusion	51

Liitteet

List of Figures

2.1	Terminal output from a Debian system when installing an Emacs package.	10
2.2	Relations between different user environments and installed programs [23].	13
3.1	The CIA-triad, a way to demonstrate conflicting security measures [36].	22
4.1	Working example of a NixOS client displaying image with Wayland and Feh.	29
4.2	Architectural graph of the test setup.	31
5.1	The QuERIES methodology is used as a reference flowchart for evaluation of security. [36]	41

List of Tables

6.1	Probabilities of each state transition occurring in a 7 state Markov chain, where s represents state, c count and p probability of the transition succeeding. C_t represents total count. Given action is "Monitor state".	45
6.2	Probabilities of observation across states.	46
6.3	Different states, defensive measures, observations and attack measures for the system.	47

1 Introduction

A Linux distribution is a bundle of the Linux kernel and a set of software products called packages [1]. A package manager is an instrument that handles building packages from either from source or pre-built binaries, resolving build-time and run-time dependencies of packages and installing, removing, and upgrading packages in user environments [1]. Every Linux device must handle its installed programs with their dependencies and configurations either imperatively or declaratively. Overwhelmingly large portion of Linux distributions fall in to the first category [2]. Both imperative and declarative systems have their strengths and weaknesses from administrative and security standpoints.

An imperative system provides updatability and modification through a destructive instrument. Popular imperative package managers are apt, apk, dnf and zypper [2]. Imperative package managers can remove and overwrite existing files which leaves the system in an inconsistent state. Different installs have by nature different states which causes many problems discussed in this thesis. As files are cross-modified through packages with package managers such as apt, upgrading can be disastrous as such systems don't support atomic rollback capabilities ¹. Due to the unpredictability, often the result can be a partially or completely broken system [2].

The reference imperative system in this thesis is Debian with its default package manager apt, due to it's popularity and relative simplicity. The reference declarative

¹Some Linux distributions using btrfs filesystem can perform a snapshot and rollback [3]

system is NixOS with it's partial namesake package manager Nix which provides declarative configuration of the whole system including the Linux kernel². Nix is configured by the Nix programming language which is inspired by purely functional languages such as Haskell. [5]

The reference architecture depicted in chapter 4 is based on NixOS as it is the most popular purely functional, declarative Linux distribution with over 100 000 packages [6]. Another good alternative would have been Guix, which has over 28 000 packages [1]. Guix has some improvements over Nix, including richer and more extensible programming environment with a Lisp-dialect configuration language, Scheme [7]. NixOS remains as the distribution of choice, as the number of packages is greater, and general support is found to be better.

This thesis focuses on the systems and information security of a reference architecture created with NixOS. Chapter 4 goes through a reference architecture of a solution that handles securely the most critical functionalities of an image displaying system.

This thesis discusses how declarative systems can be used as an improvement over imperative systems in a public information screen setting. In chapter 2 both approaches to package management is compared, and in chapters 4 and 6 a quantitative research is carried out revealing strengths and weaknesses of the reference Nix environment setup. Selected methodologies provide quantitative results which can be used to improve the security of similar declarative systems. Propositions for further research are gone through in chapter 7 and finally, chapter 8 concludes the thesis.

1.1 Research methodologies and questions

Research methodologies used in this thesis are:

²There exists an experimental project that has succeeded with BSD interoperability [4]

1. a literature review of central papers on subject themes found with prepared search statements
2. an action research using a laboratory setup
3. a quantitative research process using QuERIES methodology

Literature review will be addressed in the next section. As this thesis' main research methodology is quantitative, the gathered data points will be addressed as variables that are compared with mathematical methods. Quantitative methods are broken down in chapter 5 and 6. The central methodology is derived from QuERIES, and the information security aspect is investigated with CIA-triad. In chapter 6 section 5.4 QuERIES is compared with other metrics which are explained in subsection 1.1.2.

Quantitative methodologies are oftentimes used in conjunction with qualitative methodologies, both approaches having their strengths and weaknesses. One drawback of using qualitative methods in security framework is their inherent subjectivity. For example the Delphi technique, where a set of opinions is gathered and compared from a working group provides subjective substance for a study instead of objective perspectives [8].

The study design in this thesis is *state based*, which refers to the fact that the research methods focus on different state transitions, e.g how probable it's for an intruder to gain from partial leverage to a full control of a system. Qualitative research wouldn't alone satisfy the requirements, as investigating different state transitions without quantitative methodologies would be absurd [9].

1.1.1 Literature review

This thesis has bibliography from x sources, most of it gathered with a carefully prepared search statement. Other sources include manuals, material for research

methods and other relevant material. The search statement's results, presented in next subsection, provide good base for action research and analysis.

The literature focuses on three main concepts: embedded systems security with and without declarative components, imperative systems and NixOS as a solution. The main goal is to find literature that combines these concepts to gain platform for comparing different approaches to support the action research.

Central literature revolves around Nix and multiple texts by Eelco Dostra are cited for illustrating the nuances of a Nix ecosystem. Other declarative approaches are discussed, by Endres et. al and Van der Burg [10], [11]. These approaches also contains comparison to imperative systems, which is the central approach in chapter 2. Combining cyber security with declarative approaches were discussed by Specht et. al and Kandoi and Artke [12], [13].

Discussion from Ravi et. al and Fysarakis et. al on embedded security is discussed by [14], [15]. The concepts, however are generally too broad for this thesis' scope, so only the most fitting approaches were selected for use.

Search statement

The main search statement for this thesis is: "embedded linux" OR "declarative" AND (linux OR *nix) OR deployment OR "system update" OR (compare* AND declarative AND imperative AND system*) OR security.

The search statement was prepared to provide as relevant results as possible for this thesis. The main goal was to include the hypernoms "embedded linux", "linux" with other terms separated using the "OR" operator. The subterm (compare* AND declarative AND imperative AND system*) was chosen to broaden the search to include articles which compare declarative and interactive systems.

As security is a central theme in this thesis, the term "security" was included. Search was done on Google Scholar, and other useful material was handpicked, such

as NixOS manuals and wiki pages. Systems security and cyber security material is also included in the bibliography using search statement systems security OR cyber security. Separate search "cia-triad" and "partial observable Markov chain" AND "cybersecurity" were used to provide tangible meters for measuring cyber security. To further back up the research for comparing different metrics, term "cyber security metric methodology" was searched.

For searching specific material about embedded systems, the search statement "embedded AND security" was used. As the need for embedded toolchains was needed, statement "yocto AND buildroot" was searched. All searches were done on Google Scholar platform.

1.1.2 Research questions

The research questions for this thesis are:

1. How can a declarative system be used to improve the basic security needs of an embedded system used for displaying public media measurably?
2. What are the advantages and/or disadvantages of such system from system administrator standpoint?
3. How can a declarative system be updated from different Linux distribution securely and seamlessly?

Research question 1 is perhaps the most important and it traverses through themes of the whole thesis. The hypothesis is that traditional imperative embedded device fleets have problems that can be solved with the use of modern declarative systems. First, we aim to gain information from a specific scenario, presented in chapter 4, then in chapter 6 the gained information is analyzed and generalized as suitably as possible.

Research question 2 brings up the human element; how can a system administrator use a new palette of features adequately to provide more secure system and research question 3 handles a situation where existing system should be replaced with a NixOS system. How this could be done securely without risks and preferably easily with existing or new tooling is answered in chapter 4 section 4.4.

1.1.3 Data collection and analysis

Data collection is done with simulated red-blue team setup, where either team has a time frame where they must conduct a series of tasks. These tasks are formalized as partially observable Markov chain parameters, and analysed with QuERIES methodology. This methodology is used to gain knowledge and make the system more reliant and better with multiple iterations. Chapter 6 answers research question 1 and 2 and provides analysis for the reference system. Research question 3 is answered in chapter 4 section 4.4.

2 Declarative vs. Imperative systems

There have been different approaches to declarative modeling of systems design. Endres et al. compares declarative and imperative systems from a cloud computing standpoint, and collects systematic information on what are the strengths and weaknesses of TOSCA, IBM Bluemix, Chef, Juju, and OpenTOSCA [11]. Van der Burg and Eelco Dostra use NixOS as a solution for declaratively distributing into cloud, executing integration and system tests [10]. Most approaches researched through literature review focus on distributing to cloud. Distributing to embedded clearly remains as a niche.

Breitenbücher et al. focus on deploying into embedded and discusses the challenges an IoT user face when deploying a system. It's proven that setting up devices with mandatory scripts and other actions is a challenging task, when a number of devices should be set up. Cloud is something that is useful to be used in tandem with IoT but this thesis focuses on an *in-premises* reference solution. [16]

In this chapter, we focus on comparing different declarative approaches to the more traditional imperative models, highlighting the strengths and weaknesses of both. Specifically, examples are provided to illustrate the limitations often observed in imperative systems, particularly in terms of reproducibility, scalability and administration standpoints. Cloud-oriented approaches serve as a prime reference point

for how declarative systems can be effectively distributed and automated. I argue that similar approaches as those taken in cloud should be taken in embedded and IoT to increase security.

2.1 Imperative systems

Imperative deployment models base their functionalities through a process in which the order of events have a critical significance to the output [16]. In context of virtualization, imperative tooling can be used to form a all activities to be executed, the control flow, their execution order, and the data flow between them [11]. This kind of process is best to be used in conjunction with a formalized workflow or standard such as BPEL [11]. In contrast, declarative models don't have such specific requirements, as these models formalize the processes in the configuration files [11].

An imperative system provides updatability and modification through a destructive instrument. Popular imperative package managers, e.g can remove and overwrite existing files, which leaves the system in an inconsistent state [2]. Different installs have by nature different states, which causes many problems discussed in this thesis.

Imperative systems, while popular, have inherent problems regarding administrative traits contributing to a framework where the underlying system has **no traceability**: the implication that reproducibility is impossible, as changes to a system are not traced. Nix provides a solution for this problem with its Nix generation system. The second point contributes to the fact that with imperative systems, upgrading is more error-prone than installing from scratch. This is due to the fact that imperative systems have **unpredictable** state, from where the system should migrate to a predictable state. This causes major issues regarding upgradability. [17]

The inability to run multiple configurations side-by-side is an inherent


```
dpkg: emacs-lucid: dependency problems, but removing anyway as you requested:
emacs depends on emacs-gtk (>= 1:27.1) | emacs-lucid (>= 1:27.1) | emacs-nox (>= 1:27.1); however:
Package emacs-gtk is not installed.
Package emacs-lucid is to be removed.
Package emacs-nox is not installed.
```

Figure 2.1: Terminal output from a Debian system when installing an Emacs package.

side effect of a *stateful* system. Declarative systems don't have this problem: an arbitrary number of configurations can exist side by side, as the system is defined only by the configuration, not with the state as a component. [17]

2.1.1 Debian/Apt

An example of imperative systems' problematic nature is provided with the following demonstration. Executing shell command

```
apt install emacs
```

installs a text editor wrapped as a Linux package. The package emacs has a dependency, emacs-gtk, which can be removed with command

```
apt remove emacs-gtk
```

Another dependency, emacs-lucid can be removed with command

```
apt remove emacs-lucid
```

we can see that after removing, apt automatically installs emacs-gtk to avoid breaking the application. The package manager warns: "emacs-lucid has dependency problems, but removing anyway as you requested" as shown in figure 2.1.1. It's also noteworthy, that the manual page for apt, doesn't say anything about a possible installation side-effect of a package removal command [18]. We could forcibly remove the package by invoking

```
dpkg --remove --force-depends emacs-lucid
```

thus leaving the system in an unreliable state. Dpkg is a low-level tool associated with apt, and doesn't automatically handle dependency resolutions or further package relations [19]. What happens if we had a large number of devices, in-premises or cloud where all system commands are done imperatively? We would have a large number of devices that differ from each other, because as shown, the order of commands affect the state of the system. Time is also a factor that causes systems to diverge, as packages are not up-to-date by default. Invoking

updates the local repositories to match the download mirrors. If by technical reasons or possible user error this command is conducted in the wrong order there will be divergent systems.

Implementing a deployment model with only Debian would be a gruesome task, as the order of events which occur during the setup phase is critical. As presented by Endres et al. a formalized workflow graph would be needed to set up a reliable system. However, a Debian could be used as a host to user-space application deployment, such as Bluemix or Chef, where common DevOps practices can be used [11].

2.1.2 Gentoo/Portage

Gentoo Linux is bundled with the package manager Portage, which consists of two components: ebuild and emerge, which have similar relation as dpkg and apt. Portage is primarily source-based package manager; ebuild builds and installs packages from source and emerge resolves dependencies and handles other related issues. Portage's flagship feature is its "USE-flags", a mechanism that enables compiling source files with or without specific features. For example, a make.conf file in /etc/portage could have USE-flags specified as:

```
USE=''-X wayland''
```

and the sources wouldn't, if possible, compile with X11 support, but would with Wayland. This results in smaller binary sizes, increased performance and enhanced security through package minimalism. When there's less dependencies and installed programs, there's also less attack surface [20].

Gentoo's portage system allows sharing binaries and sources with rsync, which would be a useful feature in a server-client model, similar to architecture proposed in chapter 4, where updates are centralised [19]. An architecture with imperative components isn't proposed in this thesis, but if it would be, a suitable candidate would be something that uses Gentoo Linux as it's foundation.

One benefit from Nix is its lightweight tendency to enable system tests. Integrating system tests with a Gentoo system would require a considerable amount of work, as setting up such system needs a lot of configuration and executing commands in a correct order [21]. Gentoo definitely fits in an imperative deployment strategy but the requirement for explicit detail of every step would be error prone even for a seasoned administrator [16].

2.2 Declarative systems

Presented problems in chapter 2 can be solved using package manager that is reproducible, reliable and atomic. Package installs in Nix are in isolation from each other so that they don't have conflicting effects. This results packages being predictable and assures they work coherently even if underlying install is different. Because the packages are declared in a single set of configuration files it's trivial to reproduce the system in a different environment. The demonstrated effect in snippet 2.1.1 was a problem due to lack of isolation. When dependencies are scattered in the system instead of declared explicitly in a installed package, a faulty state could be achieved. Nix assures, that these kind of problems are out of the question. A result of this is that in a Nix system installs of same program can reside side-by-side with varying

versions [2].

As presented by Endres et al, systems can be declared, even if the underlying infrastructure is imperative by nature [11]. This thesis focuses on purely functional methodologies which fix the most prevalent issues compared with imperative models. Tools such as Chef focus on deploying on a imperative system, which causes an inherent problem with cohesion in a system that should work regardless of the underlying machine or network. Alternative deployment tools are discussed in section 2.2.1.

It's also noteworthy that many imperative package managers don't support rollback mechanisms. If the Nix configuration file is changed and the system is rebuilt with command

```
nixos-rebuild switch
```

the previous state could be recovered by

```
profile rollback
```

This is an important feature as the Nix configuration files control the whole system, they can also leave the system in an undesired state. Nix switches between *profiles*, which is a way to provide different configurations for different user environments as shown in figure 2.2 and provide atomic upgrades and rollbacks. [22]

A fundamental component of the ecosystem is Nix, a domain-specific language designed for configurations, distinguished by its functional nature and lazy evaluation. The concept of purity is central to Nix, where values remain unchanged throughout computation, and every function consistently yields the same output regardless of input [24]. The security implications of using Nix vs. an imperative system is discussed in chapter 3.

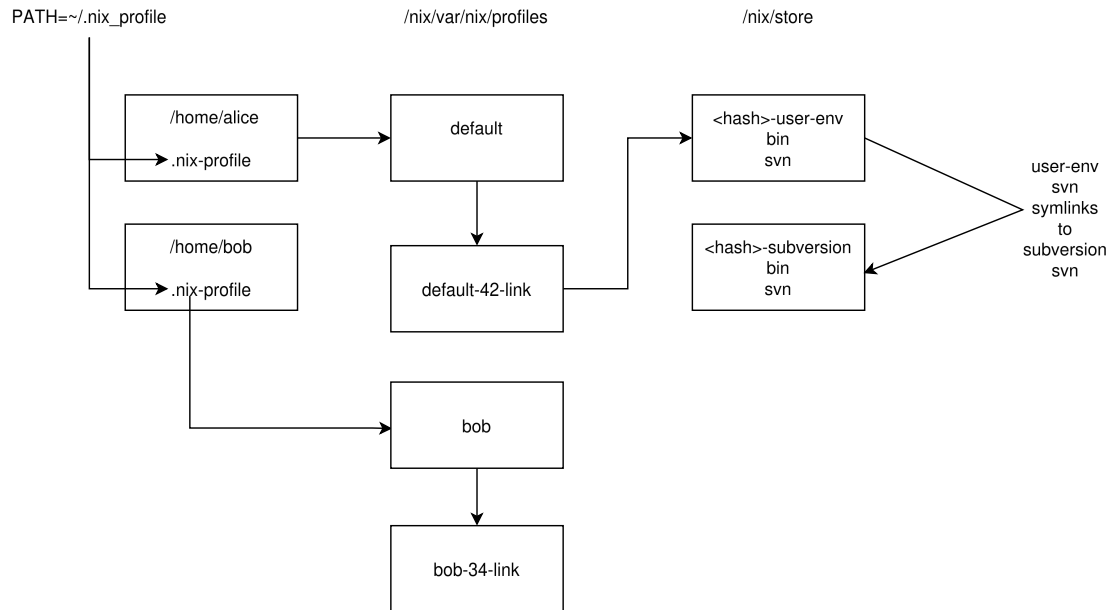


Figure 2.2: Relations between different user environments and installed programs [23].

2.2.1 Non-declarative components

Declarative distributions such as Nix can't do everything in the system in stateless manner. Some components of the system, such as databases must have a distinct state, which can't be practically declared with package manager apart from initial configurations [5]. Home directories can vary as much as the system administrator desires. For example, a configuration file for text editor vim is usually declared in the file `/home/<user>/.vimrc`. Nix provides multiple ways to perform the whole configuration process from the Nix configuration files. One way is declaring the desired `.vimrc` in the Nix configuration, as in the following snippet:

```
{
  environment.systemPackages = [
    (pkgs.vimConfigurable.customize {
      vimrcConfig.customRC = ''
        " arbitrary vim config
      '';
    })
  ]
}
```

```
];  
}
```

Nix provides also provides ways to fetch content to the system from remote URLs, and if the administrator doesn't want the system to remain "pure", they can build the system by

```
nixos-rebuild switch --impure
```

This results the system having mutable components, which can be desirable from an accessibility point of view, but can result in unpredictable behaviour if the impure components are modified. In this context this means, that the built pure components are read-only and immutable [25].

User environments (Nix profiles) can be used so that for different needs, or for different users there are multiple environments in which the user can operate as shown in figure 2.2. User environments are a successor to the concept, where installed programs either reside in `/usr/bin`, `/usr/sbin` etc. or have a symbolic link to the said directories. They can be figured as trees of symbolic links that reside also in the Nix store hence referred packages are called "activated packages". Remember, that the traditional Unix directory structure, where files are separated to `/bin/` `/sbin/` etc. doesn't completely apply. Instead, the installed programs reside usually in `/nix/store`. [2]

There are continuous build and integration services, such as Hydra, which include Nix-compatible support for handling runtime configuration and tools, such as Disnix and Charon ¹, which focus on setting up complementary infrastructure. Van Der Burg presents these new tools to replace Cfengine, Puppet and Chef, which execute operations in convergent manner, meaning that they capture what changes should be done to the machines in a specified network. [5]

These approaches have two central problems: imperative nature of handling

¹Charon is now called NixOps [26]

environment difference, and inability to guarantee configuration compatibility with a machine. Disnix is a Nix derivative that can overcome these challenges by separating logical properties from physical, and by capturing the essential aspects which form a system. The result with Disnix is a system that can be reproduced anywhere initially, and upgrading it is a trivial task. [5]

2.2.2 Home manager and flakes

NixOS environments can be build from a single `configuration.nix` file, but there are two significant configuration tools for managing NixOS systems: home manager and flakes. Home manager is an extension for managing user profiles with a declarative Nix syntax [27]. Home manager has problems with atomic rollbacks and for this reason they are not used in this thesis' examples [27].

Flakes are experimental feature of Nix, providing environments, where dependencies are pinned in a lock file, further improving reproducibility of Nix systems. A flake is defined as: "... a file-system tree whose root directory contains the Nix file specification called `flake.nix`". The usage of flakes is a good method for organising different environments within a Nix system, where it can consist of multiple flakes. Flakes, however are a experimental feature, thus out of this thesis' scope. [28]

2.2.3 Ease of updates

Updating is easy and riskless with NixOS due to atomic rollbacks. Nix handles software providing through something called "channels". Channel is a set of latest Git commits in a Nixpkgs repository, where they are divided to stable/unstable and large/small. channels. Unstable channels (large and small) have the latest commits on a rolling basis, but include less conservatively checked functionalities. Stable channels are submitted through a version number (e.g. 23.11), where a new release is published every six months. Large channels on the other hand contain a full set

of Nixpkgs binaries, when small include a subset. If a system administrator decides to submit to small channel, they have more recent updates at their disposal, but have to resort to compiling some needed packages from source. [29]

Updating a Nix system is just a manner of invoking command

```
sudo nix-channel --update
```

and, if stable release is chosen, updating the `system.stateVersion` from the `configuration.nix` file [22]. Nixpkgs is a repository of working Nix packages using a continuous integration service called Hydra. Hydra evaluates the needed Nix expression of a package, and ensures its functionality. [22]

3 Embedded system security

According to Serpanos et al. the use of embedded devices can be divided into four fields: industrial systems, nomadic environments, private spaces and public infrastructure [30]. This thesis' focus is public infrastructure, specifically information screens in a public environment. Implementing security mechanisms and policies is essential for information screens to function securely from both organizational and technical viewpoints. Implementing those policies and assuring compliance is trivial with declarative approaches with increased benefits from reliability perspective.

Embedded systems are distinct from other type of systems due to their varying nature ranging from programmable logic controllers (PLC) to larger systems, such as servers or routers. [15]. An usual embedded device conducts a specific task and possibly demands networking capabilities. Working with embedded is typically working with a limited set of resources, demanding careful design when a multitude of features are needed. Maintaining and upgrading devices to meet the continuous need of security updates. Even services such as SSH have had history of vulnerabilities, which prove that upgradability is a fundamental base of a secure system [31]. I argue that the security aspect of embedded devices could be improved significantly with the use of declarative systems as seen in chapter 2 section 2.2 and in the following section 3.2.

Embedded devices demand precision and security, as their function may be very critical for variety of safety reasons, e.g in automotive industry or healthcare ap-

plications [15], [32]. Reliability is a defining requirement for number of embedded applications; a pacemaker that doesn't function all the time reliably is completely useless. While a declarative solution itself can't fulfill all security needs, it definitely could improve the *reliability* of such systems.

As stated by Fysarakis et al. defining an access control is essential for any system to prevent unauthorized access [15]. Implementing complex access control is trivial with Nix, as the configuration files denote completely which user has accesses to which resources. Access control in a modern day embedded environment could be hard to implement in traditional imperative systems, as scaling an access control system which spans multiple devices and changing environments would require a lot of manual intervention. This is definitely one of strengths of declarative approaches: scalability is never an issue when a centralized configuration defines the systems. A declarative approach is often taken in the cloud as stated in section chapter 2 section 2.2, but implementing declarative overlays definitely needs work in the embedded field.

Implementing policies information security perspectives using a policy modeling standard, CIA-triad is discussed in section 3.3, and Nix is reflected with the use of the triads axes. Dolstra states that Nix is policy-free meaning that it contains a set of mechanisms which allow policies to be constructed with and not the other way around [33].

Embedded being a broad field, in this thesis devices are limited to those which can run Linux kernel and provide the most basic networking capabilities. These cover architectures i686, x86_64, arm64 supported by NixOS. PLCs and microcontrollers are outside of scope as NixOS needs a functional Linux kernel and a specific architecture to work.

3.1 Common embedded pitfalls

Common issues regarding embedded devices are their lack of updates, weak data integrity, and the multitude of features [15], [34]. For example, a toy teddy bear may have a audio recorder, data transfer capabilities and ability to geolocate itself. These kind of devices may lack firmware or software updates, and the data-transfer may be insecure.

A solution for secure data transfer would be TLS-encrypted messaging between clients. This could be achieved with MQTT-protocol, but configuring certificates is extra effort. Multitude of features is a definite security problem, as the user may not be aware of them at all times. In an increasing global world, importing embedded devices from unreliable sources can prove to be a security issue. The household items may or may not adhere to latest security compliance. [15]

Attack surface of embedded systems in general range from physical access to network and geolocation problems. One way of manipulating a device, apart from directly gaining access to the operating system, are side channel attacks. Analysing the power or electromagnetic properties of device input/output can be used to determine critical aspects of a device, e.g key lengths or algorithms of security measures [15], [30]. Attack surface may used to gain access, or performing denial of service attacks. Geolocating is both a privacy and security issue, as location data may be used to trace identities of device users, which can lead to e.g blackmailing, physical intrusion or other means [15]. This means that the principals of this thesis' scope could theoretically be targeted with such malicious intents.

Embedded systems have problems regarding monitoring and system administration. It's very different to have home automation system with less than 20 nodes, than to have public transport embedded fleet in a big city with 2000 nodes. As the number of devices grow, so does the challenge of monitoring and administrative tasks. Home automation has usually one person dedicated to the task: the

home owner. The hypothetical setup with 2000 devices has an exponential growth of problems. Monitoring should be trivial to automatize (e.g by using tools like Prometheus), but administrative tasks are harder to automatize, due to tasks being potentially very challenging even for dedicated system administrators. This is where Nix comes to play, as updating thousands of devices becomes trivial.

3.2 NixOS as an embedded solution

Declarative systems have advantages over imperative systems in reliability and safety aspects due to two things:

1. rollout and rollback are equally trivial tasks
2. desired configuration can be tested in a sandbox environment

The first item makes it more accessible to manage a rollout strategy, as the rollout/rollback can be done multiple times or executed completely in a replicated sandbox environment, as stated in item 2. Simpler and more straight forward practical steps give space for easier strategical planning. [13].

Kandoi et. al argue that with declarative systems, it should be nearly impossible to misconfigure in the first place the system and if faulty state is achieved, a simple rollback could undo the changes [13]. As stated in chapter 2 section 2.2, it's definitely possible to achieve faulty systems with Nix. I argue that these problems can be mitigated with a well thought rollout/rollback strategy.

Updatability is possible with many different platforms, but it's a problem when updating is a sole duty of a consumer, who may or may not have the adequate knowledge how or why they should update their systems. Lightweight updatability comes out-of-the-box with Nix, and that is something that inherently should make it more secure. Consumer products, however are out of this thesis' scope.

Nix is a double edged sword for system administration tasks. On one hand, it has a steep learning curve, but on the other hand it can make tasks that could be very challenging with traditional systems, trivial. In a well built Nix ecosystem security actions such as updating or modifying user or kernel space can be used to enhance security and in such system, any changes could easily be replicated to multiple devices, without the need for manual intervention.

Some other clear disadvantages for NixOS in embedded use is the fact that a purely functional, declarative system inherently must use disk space more than it's imperative counterparts. In the worst case scenario, if one derivation of a system takes up 1Gb of space, when making changes, the resulting system will need 2Gbs of space. The worst case scenario rarely occurs, but due to Nix's indestructive nature, this formula of disk space demands has to be considered in an embedded setting. [17]

3.3 Imperative and declarative systems from CIA-triad approach

CIA-triad can be used as a tool to show conflicts between different points of information security interests. It consists of three meters: confidentiality, integrity and availability as seen in the figure 3.1. Confidentiality can be seen as superset of privacy. Confidential data is classified with technologies such as data encryption and user privileges. Integrity means that the data has not been tampered with, and remains untouched by unauthorised parties while it's in transit or stored e.g in a server. A way of providing integrity is checking hashes of downloaded files. Availability is a user viewpoint to the accessibility of the system. When confidentiality and integrity are pushed to the extreme, availability aspect suffers, e.g when a service enforces multi-factor authentication. [35]




Figure 3.1: The CIA-triad, a way to demonstrate conflicting security measures [36].

Systems with an imperative package manager are more accessible than declarative systems as learning a new programming language with esoteric paradigm can pose extra effort. Configuring a whole Nix system demands a thorough knowledge of Nix language, and that definitely hinders the ease of access to a Nix system from a system administrator standpoint. With NixOS, an easy extent of accessibility can be achieved via planting sufficient configuration files during device setup.

Atomic systems such as Nix have great benefits towards integrity. As the "nix store", where every installation is located is read-only, it's impossible for attackers to modify the store. That's not the case, where user with root privileges can arbitrarily modify installed programs and files.

According to Nix manual, it has three main strengths in relation with security. **Security by obscurity:** combined with the unusual file system and the usage of user-environments, some malware that rely on the usual locations of installed programs may fail [37]. **Multi-user installations:** the requirement for root access nearly always widen the potential attack surface. NixOS provides a way for multiple users installing their programs through the use of user environments, hence mitigating the need for root access. This both lessens the availability aspect, as well as mitigates the programs root access. When file changes are made in user-specific

scope, a thin layer of isolation is achieved [22]. **Data integrity** is achieved both by installed programs residing in the read-only Nix store, but also them having been checked against SHA256 checksums. Moreover, the core installation resources for NixOS are GPG-signed by an administrative Nix team [37].

4 Proposed architecture solution

For gaining data points for the research, a test setup must be configured. This chapter presents a Nix ecosystem, which goes through a red–blue team testing in chapter 6. The proposed architecture solution uses a completely declarative approach introduced in section, where pros and cons of both approaches are discussed 2.2.

The proposed architecture is a referential framework for a Nix ecosystem and provide only a basis for a subset of features needed in such system. For example, remote access and tunneling would require extra work in a production setting. Note that presented solution’s scalability could be improved with the use of tools referred in section 2.2.1.

The main purpose of the architecture is to construct a system that focuses on reliability and fluid deployment tasks. Kandoi and Hartke discuss the operating large scale IoT solutions through declarative configuration APIs and conclude that it can solve multitude of scalability and extensibility challenges of IoT systems thus ensure reliable and safe operations [13]. As embedded security is the focal point of this thesis, a further analysis is presented in chapter 6.

4.1 Server

Some approaches, i.e by Van der Burg et al. provide a reference architecture with OpenSSH, Quake 3 server and Transmission services [5]. Dolstra, Vermaas and

Levy build a declarative system with simple web server [24]. This thesis' server has is moderately complex in comparison with these approaches, as it contains multiple components and a functioning server-client implementation.

A proposed architecture can be viewed in figure 4.2. In simplicity, the NixOS server runs a MQTT-broker for publishing images encoded in base64 format for the client devices. The clients are called as device fleet, and the server is called central server. The client devices submit to configured topics, and display them using a Wayland compositor, Weston.

The central server has two main functions: receiving SSH-connections for admin usage and forwarding bitmaps formatted in base64 to the clients. In production environment, the server would be the element that is in a public network, and the devices would be accessible locally (through a tunnel). SSH-connections to the machine would then be forwarded through the central server to the clients. The central server currently doesn't generate the imagery, but this could be achieved via headless browser, or other graphical tools.

MQTT is a extremely lightweight, machine-to-machine publish/subscribe protocol. It can be used on virtually every platform including microcontrollers [38]. The chosen MQTT broker and client for this project is Mosquitto. The following snippet shows how the Mosquitto server is configured in the NixOS server.

```
services.mosquitto = { enable = true; listeners = [ { users.<
  user>
    = { acl = [ "readwrite #" ]; }; settings = { cafile =
      "<ssl-path>/ca.crt"; certfile =
        "<ssl-path>/myhostname.crt"; keyfile =
          "<ssl-path>/myhostname.key"; }; } ]; };
```

The "services" statement tells us that a SystemD service is being defined. The settings section specify which certificate and key files are to be loaded to the service.

The MQTT-broker (Mosquitto in this case) publishes a message that is forwarded

to the subscribing clients via a following script.

```
nix shell nixpkgs#mosquitto --command mosquitto_pub -h localhost -t
images/test -m "$IMG_BASE64"
```

Sending could be automatized with a service using SystemD timer:

```
systemd.timers.publish-image = { timerConfig.OnCalendar =
    "*-*-*
    *: *:00"; wantedBy = [ "timers.target" ]; };
```

which invokes a specified service.

4.2 Client

Some IoT solutions prefer client's such relationship with client and server, where the client automatically searches for suitable server dynamically [13]. This thesis' client structure is static meaning it follows static addresses and forming initial connections requires manual intervention.

Both server and client are running NixOS. The client has two main functions: subscribing to media receiving and displaying the gained media which can be arbitrary. Currently, the image refreshes every second and through configuration, technically even displaying animations with this setup could be possible. Media display happens with feh, an image showing tool, that works with X server. However, this example is using Wayland, so compability layer XWayland must be used [39]. Image data messaging functions through MQTT-protocol, which is explained in the next subsection.

4.2.1 MQTT-client

The MQTT client subscribes to a topic from the following script. The image is received as base64 string, and is converted back to PNG format.

```
IP="<server ip>" TOPIC="images/test" nix shell nixpkgs#mosquitto
--command mosquitto_sub -h $IP -t $TOPIC >" <image directory path>
/image.base64" base64 -d "<image directory path>/image.base64"
>images/latest.png
```

4.2.2 Weston/XWayland

Wayland is a display protocol aiming to replace partially or fully the old X window system. Wayland functions thorough a "compositor" (server), and that provides a surface for the device to draw graphics. Wayland was selected for this project due to increased security, as the X window system has support for network transparency which broadens the attack surface. Wayland has combined server and client rendering with the Wayland compositor, so that safety-critical throughput between display server and window manager is not a concern. [39]

The example project displays an image from a directory via script:

```
sleep 5 && /nix/store/qc9j6pm6ykyx531s4kb06084mczy2l6g-feh
-3.10.1
/bin/feh -F -Z -R 1 <image-path>/latest.png
```

As the programs must be found from an absolute path, the system must generate the scripts accordingly. This is done via a SystemD service, specified as:

```

{ config, pkgs, ... }: let fehLaunch = pkgs.writeText "feh.sh"
  ''
  echo "sleep 5 && ${pkgs.feh}/bin/feh -F -Z -R 1 <image
  directory>/latest.png" > /home/user/abzug-receiver/weston/img.
  sh
  ''; initImg = pkgs.writeText "initImg.sh" '' echo
  "${pkgs.gcc}/bin/gcc <source path> img.c -o <binary path>/img"
  >
/home/user/abzug-receiver/weston/init.sh '' in {
  systemd.services."weston1" = { enable = true; unitConfig = {
    Type = "oneshot"; }; serviceConfig = { Environment =
    "XDG_RUNTIME_DIR=/var/run/user/1000"; ExecStartPre =
    ["${pkgs.bash}/bin/bash ${initImg}" "${pkgs.bash}/bin/
    bash
    <init.sh path>/init.sh" "${pkgs.bash}/bin/bash
    ${fehLaunch}" "${pkgs.bash}/bin/bash <init.sh
    path>/init.sh"]; ExecStart = "${pkgs.weston}/bin/weston
    --config=<weston configuration directory>/weston.ini";
    RestartOn = "failure"; }; wantedBy = [
    "graphical-session.target" ]; }; }

```

This wouldn't be a problem in traditional Linux distribution, but in NixOS the program locations vary from machine to machine [25]. A program resides in Nix store, with a cryptographic hash of all build inputs in its directory path. [25]. For that reason, one way of proceeding is to write a SystemD service to generate the configuration files. Note that this SystemD configuration differs in how SystemD scripts are declared usually. Most SystemD distributions have SystemD files in `/lib/systemd/system` directly or via symbolic link.

In the beginning of the configuration, variables are defined and program locations are expanded from Nix package paths. Then, in the "serviceconfig" part of the

```
[core]
idle-time=0 xwayland=true
[shell]
panel-location=""
panel-position=none**
[autolaunch]
path=<feh launcher path>/img
```



Figure 4.1: Working example of a NixOS client displaying image with Wayland and Feh.

configuration, the strings are forwarded to shell, which in part compiles sources and executes scripts. After the "ExecStartPre" section, in the "ExecStart", Weston is launched with very basic kiosk configuration:

the "[autolaunch]" only functions with compiled binaries thus the shell script is not directly executed. Instead, a program written in C is invoked, which in part invokes the shell script with parameters. The autolaunch path can handle switches, but unfortunately not parameters. With these workarounds the kiosk successfully can display media as shown in figure 4.1. Feh is launched with -R 1 parameter, which causes it to refresh the image every second. That way when a new image is uploaded, the display is also refreshed.

The sources and scripts are downloaded from Github, again via a SystemD ser-

vice. Following snippet shows the "ExecStart" part of the service.

```
ExecStart="${pkgs.git}/bin/git clone <git url> <installation path>";
```

This part of the configuration is not purely functional, as the downloaded scripts and configuration can be arbitrarily changed with correct permissions. The SystemD services impurity don't trigger the Nix language evaluation itself, so "–impure" switch isn't mandatory.

4.3 Test device

A huge benefit from declarative systems is their broad possibilities of automated system tests [21]. The phrase: "if it works on one machine, it will work on another" builds a stable foundation for such tests [22]. Different deployment strategies benefit from slightly different approaches, but as deployment strategies generally are out of this thesis' scope, only a minimal test setup is configured.

NixOS devices can be upgraded and updated through a server or locally with physical access [22]. In this thesis' example, a separate test machine is configured, and it's programs are replicated to other devices as seen on figure 4.2. Updates are done manually with command:

```
nix build --eval-store auto --store ssh-ng://remote-host
```

The test device is thus replicated to all remote hosts manually, but an automation script would be useful as the setup scales.

4.4 Installing new devices

Installing devices could made trivial, as it's trivial to generate new NixOS images from pre-existing configuration files [22]. The problem is, however, that in many cases a legacy architexture exists, which needs to be overwritten. This section




Figure 4.2: Architectural graph of the test setup.

focuses on research question 3, providing a solution for migrating from existing architectures.

If the whole setup needs to be replicated to an existing device fleet, a great help would be open source project "NixOS anywhere". NixOS anywhere specifies minimal specifications: "Unless you're using the option to boot from a NixOS installer image, or providing your own kexec image, it must be running x86_64 Linux with kexec support. Fortunately, most modern x86_64 Linux systems have kexec support. By providing your own image you can also perform kexec for other architectures e.g aarch64". NixOS anywhere also has a requirement that the devices should be available at a public network (not only wireless LAN), which wouldn't probably be a problem in production environments. NixOS anywhere also works only with NixOS flakes, so this thesis' configuration would need to be contained in a flake. [40] The installation would contain the following steps:

1. Run "curl -L https://github.com/nix-community/nixos-images/releases/download/nixos-unstable/nixos-kexec-installer-noninteractive-x86_64-linux.tar.gz | tar -xzf- -C /root /root/kexec/run" for booting separate kernel for installation.
2. Generate a minimal configuration file: "nixos-generate-config --no-fileSystems --root /mnt"
3. Add ssh-keys to the configuration
4. Upload the flake from the server to the device "nix run github:nix-community/nixos-anywhere --flake <path to configuration><configuration name> root@<ip address>"

Another way would be setting up NixOS from the installation media and setting up manually or creating a installation media with Nixos generators project (<https://github.com/nix-community/nixos-generators>). Nixos generators is avail-

able from Nixpkgs, and can thus be installed to a NixOS development machine or server. Invoking

```
nixos-generate -f iso -c <configuration.nix path>
```

would result in a ISO format image, ready for installation. [40]

5 Security standpoints

Security is inherently challenging to measure adequately, due to its complex and chaotic nature. Qualitative analysis may also result in subjective output. As such, no unambiguous standard of measuring security can be provided [8]. To overcome these challenges, several metric systems are compared, and the one that provides the most precise answers specifically for this thesis' cause is chosen.

In this chapter, selection process for meters, which in part are used for gaining quantified data, is presented. When finally quantitative metrics are gained with the help of two paradigms, GQM and specifically its superset SMART, a red-blue team layout is erected using QuERIES methodology. SMART is used in conjunction with the literature review to reveal the most suitable methodology for this thesis.

5.1 A brief history of security metrics

The history of security metrics begin from Trusted Computer System Evaluation Criteria (TCSEC), also known as the Orange Book from 1983 which popularized many terms still in use today, such as identification, authentication and authorization. [41]

While US National Bureau of Standards (NBS), the organization later formed as the National Institute of Standards and Technology (NIST) tried to standardize security, in the 1990s it became evident that a system should adhere too strongly to the definitions of the Orange Book and the follow up project, Common Criteria

to be used accordingly. Later, multiple standards became popularized, e.g System Security Engineering Capability Maturing Model (SSE-CMM), which can be used as a sort of checklist from system design from the ground up. [41]

Later it was observed that systems design is only a part of a successful security strategy, and operational practices played a bigger role than expected, something to be addressed in 1995 NIST Computer Security Handbook, which has evolved to provide ground to combat modern issues. [41]

As metrics can be used as a tool for decision making, the strategical approach of the mentioned publishes is important. It's noteworthy that the strategies (the Orange Book, SSE-CMM, etc.) begin to measure security by compliance to defined ratings. Later in 2000s more mathematical approaches were taken, which are taken into account in section 5.4. [41]

5.2 Choosing security metrics

Security is something that is challenging to measure due to it's complex nature. A GQM (Goal, question, metric) paradigm helps to choose appropriate metrics: first there must be a set goal to a organisation, then a formulated question for each goal. These answers are then reflected to gain the desired metric. This strategical approach is perhaps too broad for this thesis' scope, but aligns well with an usual organisational strategy. [42]

A more appropriate tool for this task would be SMART - a set of inputs to evaluate meter systems' suitability. These inputs describe how specific, measurable, attainable, relevant and timely the methodology is. [43]

In cyber security, being specific is very important and a common issue with security meters is that they either cover too many topics and are without precise definitions, or they are too specific to be generalized to broader scope of situations [8]. This thesis' results are important to be measurable, as the research orbits

around system states, and the research aims to measure with what outputs do the state transitions resolve to.

To be attainable is relevant to this context for the reason that a thesis has different scope than e.g a huge organization, and the proposed setup has to go through a check: are the metrics' goals achievable. Relevance has to do with risk assessment: how important it's to measure something related to it's value. Risk assessment is gone through thoroughly in 6 section 6.1. Time-bounding signifies the importance of time as a meter; a system that can be penetrated in a minute can definitely be seen as weaker than a system that takes years to be compromised.

5.3 Measuring security

In this section, different methodologies and perspectives gained through literature review for cybersecurity are discussed, and potential methodologies are compared to gain the most adequate metric system for usage through SMART process [43].

Security metrics can be divided to address four separate themes: **System vulnerabilities**; measuring vulnerabilities can be applied to user, interface-induced, password, and software vulnerabilities. Users are always susceptible to e.g phishing attacks or malware infection, where a user of an arbitrary system is the definitive attack vector. Interface-induced vulnerabilities refer to attack vectors related to open ports and endpoints. [44]

Password vulnerabilities refer to situations where password can be computationally cracked. This is relatively simple to measure, as it can be estimated how much time it takes to crack a password, or with the use of statistical password guessability. Software vulnerability on the other hand is a very usual way for a cyber breach to take place. This kind of vulnerabilities can be measured, thus also estimated with the help of exploitations in the past. Time is the essential element here, as the time to patch a software vulnerability is a central metric. [44]

Defense measures can be applied to strength of reactive, preventive proactive and overall defenses. Reactive measures include blacklisting, a lightweight mechanism to prevent e.g a botnet to harm the protected system by blacklisting IP-addresses related to the botnet. For measuring defence, the reaction time is essential, and most importantly, the gained meter to measure preventive defense. Blacklisting can also be used as preventive and proactive measure, as a pre-filled blacklist can be used with desired parameters. [9], [44]

Overall defenses can be measured with the combination of all defensive measures and with the use of penetration testing in a red-blue team setup. Penetration testing aims to gain a result, also known as penetration resistance, which is a meter, indicating cost or time that the red team must spend in case of a successful system compromisal. [9], [44]

Threats: zero-day vulnerabilities can be measured from two perspectives: lifetime of zero-day vulnerability and the number of nodes that are compromised as a result. Malware spreading can be traced with the parameter infection rate, which is defined as infected node per a time unit. Attack evasion is measured using either obfuscation prevalence metric, or structural complexity metric which provide information on obfuscating gained samples e.g by encrypting, or the target system's complexity measured by runtime. [9], [44]

Situations, which can relate to security state, security incidents and security investments. Security state has multiple parameters, including incident rate and blocking rate. Security investments on the other hand measure the budget percentage funneled towards security, and the return of such investment. [44]

5.4 Methodologies in comparison

Today, there exists cybersecurity metrics based on quantified mathematical models, which are prevalent for this thesis. Three different methodologies are discussed, and

one is picked for measuring the security of this thesis' architecture implementation. SMART is used to choose the metrics, and the fact that this thesis' architecture implementation is state-based. All the following metric systems are **measurable**, but some fit better especially according to **time-related** and **relevance** axes.

The literature review provided three central methodologies from variably different perspectives. Complex mathematical models are presented by Alshammari et al. are too broad [45]. This thesis' scope is limited, and this methodology would fit better a wider cyber security setting.

A methodology based on object-oriented thinking, followed by UML-graphs is adequate in many contexts, but as stated in the paper: "This measurement is a comparative one. It can be used to compare various alternative designs of the same class with respect to their security properties". As it has been stated, this thesis focus isn't comparative, as all the critical comparison has already done in chapter 2.

Hidden Markov models presented by Wang et al. are close what is the end goal of the research is in this thesis [46]. The **time-related** aspect would be satisfactory, as the hidden Markov process deals with a time parameter. The problem, however is regarding the generality of the methodology, and something more **specific** would be a better fit for this thesis. **Relevance** also is an issue, as using presented Hidden Markov models would be mathematically challenging, and perhaps too demanding for the scope of this thesis.

The last and the most fitting methodology would be one presented in papers by Carin et al and Hughes, Jeff and Cybenko [47] [36]. The QuERIES methodology is delved deeper in section 5.1, and it's **time-related**, **relevance** and **specific** axes are a near-perfect match for our goals as the model itself is relatively simple and provides shifting probabilities from states, which serves this thesis' study design well.

5.5 Quantitative metrics

Selecting carefully a metrics system includes asserting our goals and questions. Our goal is to discover this thesis' architecture proposals tenacity in a simulated setting. The main goals reside in this thesis' two research questions:

- How can a declarative system be used to improve the basic security needs of an embedded system used for displaying public media?
- What are the advantages and/or disadvantages of such system from system administrator standpoint?

The proposed architecture solution presented in chapter 4 will go through a red and blue team inspection, complying with the QuERIES model 6.

5.5.1 QuERIES

QuERIES model consists of number of steps that

1. model the problem - by conducting a risk assessment of the attack surface and the value of the possible intrusion
2. model the possible attacks - build an attack graph of intruding through vulnerabilities or other means
3. quantify the models - by conducting a controlled red team attack and provide quantified results for the said attack
4. use the results - use blue team methodologies to provide increased protection against the exposed problems

First, risks are assessed of the attack surface due as a blue team task. It's very important for blue team to know what are the most critical points of the

attack surface, and it's also used as the base for quantitative analysis. Value of the intrusion can also be used for the reward model for analysis. As seen in figure 5.1, the methodology is applied *iteratively*, i.e the steps are repeated as many times as needed for the system to be secure.

Modeling the possible attacks is a task for the red team – by constructing an attack graph, the opposing forces have a plan, which can be used as a template for analysis. By documenting all steps, we gain academical ground for conducting a thorough analysis.

In this thesis, models are quantified with the use of time framing. Both teams have limited amount of time to conduct their tasks, and probability for succeeding a certain task is calculated with formula

$$\frac{t_e}{t_t}$$

where t_e stands for elapsed time and t_t for maximum time that can be used which is the same for all tasks.

5.5.2 QuERIES as a central methodology

QuERIES draws inspiration from computer science, game theory, control theory, and economics, thus is a complex answer to a complex question. It is stated that it can be used as an alternative to popular methodologies such as red teaming or black-hat analysis used commonly in risk-assessment. [47]

QuERIES is proposed to have potentially significant usage in DoD (Department of Defense) and in private sector [47]. Initial testing of QuERIES in small-scale, realistic scenarios presented by Carin et al. suggest that the methodology can in fact be used as to improve risk-assessment more complex settings [47]. This thesis follows similar steps: first the QuERIES methodology is used to assess risks and then they are generalized with strict constraints in mind.




Figure 5.1: The QuERIES methodology is used as a reference flowchart for evaluation of security. [36]

As stated by Hughes and Cybenko the result of QuERIES isn't binary: the attacker must think about the most optimal timeframe to stop the operation [36]. The strategy uses **open** and **closed** loop decision algorithm for deciding when to stop trying. Closed loop decision algorithm constantly evaluates when is the optimal time to stop trying and open loop means that the system has pre-defined goals for evaluation [47]. As a positive side-effect by gaining the probabilities through red-team evaluation, the system is thoroughly tested and improved. This is a valuable metric, providing insight on how does the value of the system transition in the worth of breaching it in certain time, thus reflecting the true cost of the attack. As stated by Hughes and Cybenko, if the value is high enough related to the value of the gained value, they will less probably perform the attacks [36].

5.5.3 Partially observable Markov decision process

Lastly, using the results of the POMDP can be used to increase the protection against the discovered problems. POMDP provides us data on two things: found security breaches and their security implications through the reward function [47]. The models and attacks are quantified through a partially observable Markov decision process (POMDP) which contains these seven steps:

1. Define possible states the system can be in
2. Define the actions the system can take
3. Define the possible observations the system can take
4. Define the transition probabilities of the system
5. Define the observation probabilities of the system
6. Rewards: guide the system towards the desirable actions and states.

[36]

A POMDP is used widely in this kind of applications, as both blue and red team have only partial observations related to the system [48]. The blue team can't be completely certain that the system is secure and the red team cannot perform fully reliably as systems and environments differ from each other. As mentioned in the previous section, the focus of the QuERIES analysis is the time to reverse-engineer the system, thus emphasizing the importance of only partial observations.

6 Analysis

As mentioned in chapter 5 section 5.1, the QuERIES inspired methodology is used in this chapters analysis. Tables 6.3 and 6.2 show relevant information about the calculations. In table 6.3 is listed all the possible states (S), defender actions (D), attacker actions (A) and observations (O). Then, every transition from state to another state was calculated as a probability, shown in table 6.2. The test counts are seen from the table cells, with their corresponding probability. Note that every row adds up to 1.0.

6.1 Modeling the problem and quantifying the models

The example project of this thesis is an image showing system that could be used e.g for advertisement, public transport timetables or practically anywhere where static media should be presented. Our use case focuses on displaying arbitrary media in a public location. The results, if an intruder should gain unauthorised access, would be anywhere from displaying improper imagery, to succeeding in displaying propaganda or other unwanted content. Unauthorised access could have a negative economic effect for the service provider, as every organisation displaying media want to remain credible among users.

The attack surface of the example project focuses on physical access and vulner-

s	c_t	s_0 c,p	s_1 c,p	s_2 c,p	s_3 c,p	s_4 c,p	s_5 c,p	s_6 c,p	s_7 c,p
s_0									
s_1									
s_2									
s_3									
s_4									
s_5									
s_6									
s_7									

Table 6.1: Probabilities of each state transition occurring in a 7 state Markov chain, where s represents state, c count and p probability of the transition succeeding. C_t represents total count. Given action is "Monitor state".

abilities in remote connections. With MQTT-messaging, SSH and display protocols, internal and external messaging happens.

If a unauthorised access would happen, the results would probably affect a part of the society, as the arbitrary content could gain media and social media attention. This public humiliation would definitely affect the credibility of the service provider, as well as the customer. Possible propaganda could affect society by spreading false information, or possibly bringing up societal issues via activism. Either way, this would be unwanted from the perspective of the service provider, customer and users.

6.2 Modeling the possible attacks

In table 6.3, the first column describes states the system can be in. The second and the third column state defensive actions, and observations of the system. The fourth column contains template of the attacks that could be conducted. After defining the starting layout, probabilities are calculated based on gathered empirical evidence.

The research setup is a simulated targeted red–blue scenario. For calculating the reward values R script with library "pomdp" was used.

o	c_t	s_0	s_1	s_2	s_3	s_4	s_5	s_6	s_7
		c,p	c,p	c,p	c,p	c,p	c,p	c,p	c,p
O_0									
O_1									
O_2									
O_3									
O_4									

Table 6.2: Probabilities of observation across states.

S0-S7	D0-D4	O0-O4	A0-A4
Idle	Monitor system	Normal operation	Intercept MQTT messaging
Receive media through MQTT	Patch system	Detected suspicious activity	Compromise Github repository
Set up SystemD services	Shutdown system	Detected system error	Gain physical access to device
Start Weston	Isolate device	Detected unusual media display	Exploit vulnerabilities in display
Display media			Exploit vulnerabilities in SSH connections
Error state			
Partial loss of system			
Complete loss of system			

Table 6.3: Different states, defensive measures, observations and attack measures for the system.

A reward function was calculated using the mentioned R program, and the output was x . A weight of 10 was used for positive results, and -100, if something was to be compromised. This weight distribution is due to the fact that even if blue team succeeds most of time, the results of failure are much worse than a succeeding result from the blue team [47]. The discount constant is used in calculations as shown in A, where it influences the priority of immediate versus future rewards [48]. Our case signifies the importance of both, so value of 0.75 was used.

6.3 Using the results

The results are investigated through the probabilities, which are represented in the table 6.2. The reward score is taken in to account on how successful/unsuccessful the setup is from a security perspective.

6.3.1 Red team implications

6.3.2 Blue team implications

[42]

7 Further research

One big limitation of NixOS is the fact that it mainly supports only x86_64, i686 and arm64 platforms [22]. This is not generally enough, as many different architectures are used in embedded, thus limiting the use-cases of NixOS. [15].

Buildroot and Yocto are toolchains aiming to produce bootable Linux environments. In other words, these toolchains can be used to create own Linux distributions. A base system and set of tooling could be used as a cross-compiled platform, where a custom package manager, which would either have cross-compiled binaries within its reach, would configure the user space.

Building custom images for is usually done with either Buildroot or Yocto. Both have their strengths and weaknesses, but in depth analysis of embedded Linux toolchains is outside of this thesis' scope. Instead, this chapter focuses on theoretical setup, where root filesystem and base systems can be created. In practice, both Buildroot and Yocto could perform this task. [49]

7.1 Further research questions

This chapter aims to answer research question 4, listed in chapter 1.1.2. Definitely, a declarative approach could be used with other processor architectures. This would require work on a declarative package manager, which would be used in conjunction with base systems created with Buildroot or Yocto.

Questions for further research include:

1. What set of tools would be optimal in creating a true cross-platform declarative package manager (e.g programming languages, possible serialization languages etc.)
2. How could the new system handle common system security and administration problems better?
3. What kind of ecosystem would be possible to create with a declarative embedded fleet regarding client/server or publish/subscribe models?

These questions provide base for further research, which would be needed, as currently no such system exists. A true cross-platform, purely functional, declarative system with atomic rollbacks would be a fresh newcomer into the world of embedded Linux.

8 Conclusion

Nix can be a useful base for issuing a purely functional, declarable systems while providing an advanced rollback mechanism. Configuring it depending on the system administrator may be gruesome, but it can dodge the usual pitfalls of more popular systems. While the usage of Nix language demands proficiency, it is a solution to distribute, update and administrate more secure systems. NixOS handles exceptionally easily features such as Wayland and audio driver setup. In the end, NixOS is very predictable and the statement "if it works on one machine, it works on another" proves to be true.

This thesis' problems could be solved with many different solutions, but the Nix approach is rather unique as it makes reproducible systems very straight forward. As the system is defined in the `configuration.nix` file, there's little that could go wrong even when changing hardware. The same couldn't be said of mentioned Gentoo or Debian Linux distributions.

As far as security implications are concerned, Nix is a good tool for asserting compliance. In purely functional environment, anything that needs compliance can be checked from the configuration files. As disk encryption is trivial and read-only Nix store provides increased integrity, a more secure system can be achieved also through the actions of a system administrator. A system administrator must adhere at to Nix philosophy with the Nix language and as all the configured parts can be viewed quickly. For example, a question "what ports are in use in each client

device” can be answered very quickly and in precise manner by just viewing the configuration files.

As mentioned, popular embedded Linux distribution creating tools, Buildroot and especially Yocto have also steep learning curves. They also lack the ease of updatability provided by Nix, which is a significant problem regarding security.

The research proved that a purely functional Nix can be a time-consuming task to configure correctly. While NixOS is a niche Linux distribution, and probably remains as such, the concepts, however I assume will live on in future. A purely functional declarative system with atomic rollbacks is something I’ll await to be iterated in the future with further user availability in mind.

Lähdeluettelo

- [1] G. G. contributors, *Packages & 2014; GNU Guix* — *packages.guix.gnu.org*, <https://packages.guix.gnu.org/>, [Accessed 24-09-2024].
- [2] E. Dolstra and A. Löb, “Nixos: A purely functional linux distribution”, in *Proceedings of the 13th ACM SIGPLAN international conference on Functional programming*, 2008, pp. 367–378.
- [3] *System Recovery and Snapshot Management with Snapper | Reference | open-SUSE Leap 15.0* — *doc.opensuse.org*, <https://doc.opensuse.org/documentation/leap/archive/15.0/reference/html/book.opensuse.reference/cha.snapper.html>, [Accessed 16-09-2024].
- [4] *GitHub - nixos-bsd/nixbsd: An unofficial NixOS fork with a FreeBSD kernel* — *github.com*, <https://github.com/nixos-bsd/nixbsd>, [Accessed 09-11-2024].
- [5] S. van der Burg, *A Reference Architecture for Distributed Software Deployment*. Citeseer, 2013.
- [6] *NixOS Search* — *search.nixos.org*, <https://search.nixos.org/packages>, [Accessed 24-09-2024].
- [7] L. Courtès, “Déploiements reproductibles dans le temps avec gnu guix”, *GNU/Linux Magazine*, 2021.
- [8] A. J. A. Wang, “Information security models and metrics”, in *Proceedings of the 43rd annual Southeast regional conference-Volume 2*, 2005, pp. 178–184.

- [9] A. Ramos, M. Lazar, R. Holanda Filho, and J. J. Rodrigues, “Model-based quantitative network security metrics: A survey”, *IEEE Communications Surveys & Tutorials*, vol. 19, no. 4, pp. 2704–2734, 2017.
- [10] S. van der Burg and E. Dolstra, “Declarative testing and deployment of distributed systems”, *Technical Report Series TUD-SERG-2006-020*, 2010.
- [11] C. Endres, U. Breitenbücher, M. Falkenthal, O. Kopp, F. Leymann, and J. Wettinger, “Declarative vs. imperative: Two modeling patterns for the automated deployment of applications”, in *Proceedings of the 9th International Conference on Pervasive Patterns and Applications*, Xpert Publishing Services (XPS), 2017, pp. 22–27.
- [12] E. Specht, R. M. Redin, L. Carro, L. d. C. Lamb, E. F. Cota, and F. R. Wagner, “Analysis of the use of declarative languages for enhanced embedded system software development”, in *Proceedings of the 20th annual conference on Integrated circuits and systems design*, 2007, pp. 324–329.
- [13] R. Kandoi and K. Hartke, “Operating large-scale iot systems through declarative configuration apis”, in *Proceedings of the 2021 Workshop on Descriptive Approaches to IoT Security, Network, and Application Configuration*, 2021, pp. 22–25.
- [14] S. Ravi, A. Raghunathan, P. Kocher, and S. Hattangady, “Security in embedded systems: Design challenges”, *ACM Transactions on Embedded Computing Systems (TECS)*, vol. 3, no. 3, pp. 461–491, 2004.
- [15] K. Fysarakis, G. Hatzivasilis, K. Rantos, A. Papanikolaou, and C. Manifavas, “Embedded systems security challenges”, in *Measurable security for Embedded Computing and Communication Systems*, SCITEPRESS, vol. 2, 2014, pp. 255–266.

- [16] U. Breitenbücher, K. Képes, F. Leymann, and M. Wurster, “Declarative vs. imperative: How to model the automated deployment of iot applications?”, *Proceedings of the 11th advanced summer school on service oriented computing*, pp. 18–27, 2017.
- [17] E. Dolstra and A. Hemel, “Purely functional system configuration management.”, in *HotOS*, 2007.
- [18] Canonical, *Ubuntu Manpage: Apt - command-line interface — manpages.ubuntu.com*, <https://manpages.ubuntu.com/manpages/bionic/man8/apt.8.html>, [Accessed 22-05-2024].
- [19] G. K. Thiruvathukal, “Gentoo linux: The next generation of linux”, *Computing in science & engineering*, vol. 6, no. 5, pp. 66–74, 2004.
- [20] L. Wang, S. Jajodia, and A. Singhal, *Network security metrics*. Springer, 2017.
- [21] S. Van Der Burg and E. Dolstra, “Automating system tests using declarative virtual machines”, in *2010 IEEE 21st International Symposium on Software Reliability Engineering*, IEEE, 2010, pp. 181–190.
- [22] *NixOS Manual — nixos.org*, <https://nixos.org/manual/nixos/stable/>, [Accessed 22-05-2024].
- [23] *User Environment - NixOS Wiki — nixos.wiki*, https://nixos.wiki/wiki/User_Environment, [Accessed 09-11-2024].
- [24] E. Dolstra, R. Vermaas, and S. Levy, “Charon: Declarative provisioning and deployment”, in *2013 1st International Workshop on Release Engineering (RELENG)*, IEEE, 2013, pp. 17–20.
- [25] E. Dolstra, A. Löb, and N. Pierron, “Nixos: A purely functional linux distribution”, *Journal of Functional Programming*, vol. 20, no. 5-6, pp. 577–615, 2010.

- [26] *Nix/Nixpkgs/NixOS* — *github.com*, <https://github.com/nixos>, [Accessed 08-11-2024].
- [27] *Home Manager Manual* — *nix-community.github.io*, <https://nix-community.github.io/home-manager/>, [Accessed 24-05-2024].
- [28] *Flakes - NixOS Wiki* — *nixos.wiki*, <https://nixos.wiki/wiki/Flakes>, [Accessed 22-05-2024].
- [29] *Nix channels - NixOS Wiki* — *nixos.wiki*, https://nixos.wiki/wiki/Nix_channels, [Accessed 22-05-2024].
- [30] D. N. Serpanos and A. G. Voyiatzis, “Security challenges in embedded systems”, *ACM Transactions on embedded computing systems (TECS)*, vol. 12, no. 1s, pp. 1–10, 2013.
- [31] *History of SSH / SecOps® Solution* — *secopsolution.com*, <https://www.secopsolution.com/blog/history-of-ssh>, [Accessed 18-03-2024].
- [32] N. Turab and Q. Kharma, “Secure medical internet of things framework based on parkerian hexad model”, *International Journal of Advanced Computer Science and Applications*, vol. 10, no. 6, 2019.
- [33] E. Dolstra, M. De Jonge, E. Visser, *et al.*, “Nix: A safe and policy-free system for software deployment.”, in *LISA*, vol. 4, 2004, pp. 79–92.
- [34] R. A. Kemmerer, “Cybersecurity”, in *25th International Conference on Software Engineering, 2003. Proceedings.*, IEEE, 2003, pp. 705–715.
- [35] G. Pender-Bey, “The parkerian hexad”, *Information Security Program at Lewis University*, 2019.
- [36] J. Hughes and G. Cybenko, “Quantitative metrics and risk assessment: The three tenets model of cybersecurity”, *Technology Innovation Management Review*, vol. 3, no. 8, 2013.

- [37] *Security - NixOS Wiki* — *nixos.wiki*, <https://nixos.wiki/wiki/Security>, [Accessed 22-05-2024].
- [38] *MQTT Version 3.1.1* — *docs.oasis-open.org*, <http://docs.oasis-open.org/mqtt/mqtt/v3.1.1/os/mqtt-v3.1.1-os.html>, [Accessed 22-05-2024].
- [39] *Wayland* — *wayland.freedesktop.org*, <https://wayland.freedesktop.org/docs/html/>, [Accessed 18-09-2024].
- [40] *GitHub - nix-community/nixos-anywhere: Install nixos everywhere via ssh [maintainer=@numtide]* — *github.com*, <https://github.com/nix-community/nixos-anywhere>, [Accessed 01-10-2024].
- [41] J. Bayuk and A. Mostashari, “Measuring systems security”, *Systems Engineering*, vol. 16, no. 1, pp. 1–14, 2013.
- [42] Y. V. Papazov, “Cybersecurity metrics”, *NATO Science and Technology Organization (STO)*, pp. 1–18, 2019.
- [43] S. C. Payne, “A guide to security metrics”, *SANS Institute Information Security Reading Room*, 2006.
- [44] M. Pendleton, R. Garcia-Lebron, J.-H. Cho, and S. Xu, “A survey on systems security metrics”, *ACM Computing Surveys (CSUR)*, vol. 49, no. 4, pp. 1–35, 2016.
- [45] B. Alshammari, C. Fidge, and D. Corney, “Security metrics for object-oriented class designs”, in *2009 Ninth International Conference on Quality Software*, IEEE, 2009, pp. 11–20.
- [46] H. Wang, S. Roy, A. Das, and S. Paul, “A framework for security quantification of networked machines”, in *2010 Second International Conference on COMmunication Systems and NETworks (COMSNETS 2010)*, IEEE, 2010, pp. 1–8.

-
- [47] L. Carin, G. Cybenko, and J. Hughes, “Cybersecurity strategies: The queries methodology”, *Computer*, vol. 41, no. 8, pp. 20–26, 2008.
- [48] J. C. M. Ashely S. M McAbee Murali Tummala, *The use of partially observable Markov decision processes to optimally implement moving target defense*, <https://scholarspace.manoa.hawaii.edu/server/api/core/bitstreams/86f32196-380c-443a-88b5-07f07649137e/content>, [Accessed 23-05-2024].
- [49] F. Vasquez and C. Simmonds, *Mastering Embedded Linux Programming: Create Fast and Reliable Embedded Solutions with Linux 5.4 and the Yocto Project 3.1 (Dunfell)*. Packt Publishing Ltd, 2021.

Appendix A Partially observable Markov decision process with R

The following R code calculates the result score from the study in 6.

```
library("pomdp")
```

```
CYBSEC_POMDP <- POMDP(  
  states = c("Idle", "Receiving media through MQTT", "Setting up SystemD service",  
  actions = c("Monitor system", "Reboot system", "Update system", "Display media",  
  observations = c("Normal operation", "Detected suspicious activity", "Detected suspicious activity",  
  transition_prob = list(  
    "Monitor system" = "uniform",  
    "Reboot system" = "uniform",  
    "Update system" = "uniform",  
    "Display media" = "uniform",  
    "Recover from error" = "uniform",  
    "Intercept MQTT messaging" = "uniform",  
    "Compromise Github repository" = "uniform",  
    "Gain physical access to the device" = "uniform",  
    "Exploit vulnerabilities in display" = "uniform",  
    "Exploit vulnerabilities in SSH" = "uniform"
```

```
),  
observation_prob = list(  
  
  "Monitor system" = "uniform",  
  "Reboot system" = "uniform",  
  "Update system" = "uniform",  
  "Display media" = "uniform",  
  "Recover from error" = "uniform",  
  "Intercept MQTT messaging" = "uniform",  
  "Compromise Github repository" = "uniform",  
  "Gain physical access to the device" = "uniform",  
  "Exploit vulnerabilities in display" = "uniform",  
  "Exploit vulnerabilities in SSH" = "uniform"  
  #"Normal operation" = "uniform",  
  #"Detected suspicious activity" = "uniform",  
  #"Detected system error " = matrix(c(0.85, 0.15, 0.15, 0.85), nrow = 2,  
  #"Detected unusual media display" = matrix(c(0.85, 0.15, 0.15, 0.85), nrow = 2,  
  #"Detected unauthorised access" = matrix(c(0.85, 0.15, 0.15, 0.85), nrow = 2,  
),  
reward = rbind(  
  R_("Monitor system", "Idle", "*", "*", 10),  
  R_("Monitor system", "Receiving media through MQTT", "*", "*", 10),  
  R_("Monitor system", "Setting up SystemD services", "*", "*", 10),  
  R_("Monitor system", "Starting Weston", "*", "*", 10),  
  R_("Monitor system", "Displaying media", "*", "*", 10),  
  R_("Monitor system", "Error state", "*", "*", 10),  
  R_("Monitor system", "Partial loss of system", "*", "*", 10),
```

```
R_("Monitor system", "Complete loss of system", "*", "*", 10),

R_("Reboot system", "Idle", "*", "*", 10),
R_("Reboot system", "Receiving media through MQTT", "*", "*", 10),
R_("Reboot system", "Setting up SystemD services", "*", "*", 10),
R_("Reboot system", "Starting Weston", "*", "*", 10),
R_("Reboot system", "Displaying media", "*", "*", 10),
R_("Reboot system", "Error state", "*", "*", 10),
R_("Reboot system", "Partial loss of system", "*", "*", 10),
R_("Reboot system", "Complete loss of system", "*", "*", 10),

R_("Update system", "Idle", "*", "*", 10),
R_("Update system", "Receiving media through MQTT", "*", "*", 10),
R_("Update system", "Setting up SystemD services", "*", "*", 10),
R_("Update system", "Starting Weston", "*", "*", 10),
R_("Update system", "Displaying media", "*", "*", 10),
R_("Update system", "Error state", "*", "*", 10),
R_("Update system", "Partial loss of system", "*", "*", 10),
R_("Update system", "Complete loss of system", "*", "*", 10),

R_("Display media", "Idle", "*", "*", 10),
R_("Display media", "Receiving media through MQTT", "*", "*", 10),
R_("Display media", "Setting up SystemD services", "*", "*", 10),
R_("Display media", "Starting Weston", "*", "*", 10),
R_("Display media", "Displaying media", "*", "*", 10),
R_("Display media", "Error state", "*", "*", 10),
R_("Display media", "Partial loss of system", "*", "*", 10),
```

R_("Display media","Complete loss of system", "*", "*", 10),

R_("Recover from error", "Idle", "*", "*", 10),

R_("Recover from error", "Receiving media through MQTT", "*", "*", 10),

R_("Recover from error", "Setting up SystemD services", "*", "*", 10),

R_("Recover from error", "Starting Weston", "*", "*", 10),

R_("Recover from error", "Displaying media", "*", "*", 10),

R_("Recover from error", "Error state", "*", "*", 10),

R_("Recover from error", "Partial loss of system", "*", "*", 10),

R_("Recover from error", "Complete loss of system", "*", "*", 10),

R_("Intercept MQTT messaging", "Idle", "*", "*", -100),

R_("Intercept MQTT messaging","Receiving media through MQTT", "*", "*", -100),

R_("Intercept MQTT messaging","Setting up SystemD services", "*", "*", -100),

R_("Intercept MQTT messaging","Starting Weston", "*", "*", -100),

R_("Intercept MQTT messaging","Displaying media", "*", "*", -100),

R_("Intercept MQTT messaging","Error state", "*", "*", -100),

R_("Intercept MQTT messaging","Partial loss of system", "*", "*", -100),

R_("Intercept MQTT messaging","Complete loss of system", "*", "*", -100),

R_("Compromise Github repository", "Idle", "*", "*", -100),

R_("Compromise Github repository", "Receiving media through MQTT", "*", "*", -100),

R_("Compromise Github repository", "Setting up SystemD services", "*", "*", -100),

R_("Compromise Github repository", "Starting Weston", "*", "*", -100),

R_("Compromise Github repository", "Displaying media", "*", "*", -100),

R_("Compromise Github repository", "Error state", "*", "*", -100),

R_("Compromise Github repository", "Partial loss of system", "*", "*", -100),

```
R_("Compromise Github repository", "Complete loss of system", "*", "*"),

R_("Gain physical access to the device", "Idle", "*", "*", -100),
R_("Gain physical access to the device", "Receiving media through MQTT", "*", "*", -100),
R_("Gain physical access to the device", "Setting up SystemD services", "*", "*", -100),
R_("Gain physical access to the device", "Starting Weston", "*", "*", -100),
R_("Gain physical access to the device", "Displaying media", "*", "*", -100),
R_("Gain physical access to the device", "Error state", "*", "*", -100),
R_("Gain physical access to the device", "Partial loss of system", "*", "*"),
R_("Gain physical access to the device", "Complete loss of system", "*", "*"),

R_("Exploit vulnerabilities in display", "Idle", "*", "*", -100),
R_("Exploit vulnerabilities in display", "Receiving media through MQTT", "*", "*", -100),
R_("Exploit vulnerabilities in display", "Setting up SystemD services", "*", "*", -100),
R_("Exploit vulnerabilities in display", "Starting Weston", "*", "*", -100),
R_("Exploit vulnerabilities in display", "Displaying media", "*", "*", -100),
R_("Exploit vulnerabilities in display", "Error state", "*", "*", -100),
R_("Exploit vulnerabilities in display", "Partial loss of system", "*", "*"),
R_("Exploit vulnerabilities in display", "Complete loss of system", "*", "*"),

R_("Exploit vulnerabilities in SSH", "Idle", "*", "*", -100),
R_("Exploit vulnerabilities in SSH", "Receiving media through MQTT", "*", "*", -100),
R_("Exploit vulnerabilities in SSH", "Setting up SystemD services", "*", "*", -100),
R_("Exploit vulnerabilities in SSH", "Starting Weston", "*", "*", -100),
R_("Exploit vulnerabilities in SSH", "Displaying media", "*", "*", -100),
R_("Exploit vulnerabilities in SSH", "Error state", "*", "*", -100),
R_("Exploit vulnerabilities in SSH", "Partial loss of system", "*", "*")
```

```
      R_("Exploit vulnerabilities in SSH","Complete loss of system", "*", "*"),
    ),
    start = "uniform",
    discount = 1.0,
    name = "CYBSEC Problem POMDP"
  )
sol <- solve_POMDP(CYBSEC_POMDP)
sol
```